

Academic Universe: An AI-Powered Holistic Student Growth Intelligence Ecosystem

Author: Aashish Rajput et al.

Affiliation: Sharda University / Academic Universe Research Group

Target Venue: IEEE Transactions on Learning Technologies / ACM Conference on Human Factors in Computing Systems (CHI) / Educational Data Mining (EDM)

Manuscript Version: Version 1.0 (Feature-Frozen Architecture & Validated Manuscript Draft)

Abstract

Modern higher education systems struggle to capture, quantify, and nurture student holistic growth due to fragmented academic records, unverified co-curricular artifacts, and the lack of explainable intelligence models. This paper presents **Academic Universe**, an AI-powered, multi-layered student growth intelligence ecosystem. Academic Universe introduces a three-tiered modular architecture: an **Evidence Intelligence Layer** that standardizes heterogeneous student artifacts (academic transcripts, code repositories, problem-solving logs, and verified credentials), a validated **Skill Intelligence Engine (SIE-1.0)** that calculates deterministic proficiency ($S \in [1, 100]$) and independent confidence ratings ($C \in [0.15, 0.99]$), and a novel **Growth Intelligence Engine (GIE)** that models dynamic skill acquisition velocity (μv), projects longitudinal growth trends, and infers multi-dimensional holistic development. By integrating verified course transcripts from the Academic Universe Document Intelligence Layer (AUDIC) with empirical GitHub and LeetCode activity, the ecosystem bridges the gap between institutional academic performance and industry career readiness. We present mathematical formulations, proof-of-concept synthetic simulation benchmarks across diverse student profile archetypes, decay sensitivity analysis, and end-to-end case study scenarios demonstrating how Academic Universe provides transparent, explainable insights for students, academic advisors, and recruiters.

Keywords: Educational Data Mining, Student Growth Intelligence, Multi-Source Evidence Normalization, Skill Ontology, Explainable AI in Education, Learning Analytics.

1. Introduction

1.1 Student Growth Challenges in Modern Higher Education

Higher education institutions face an unprecedented challenge in tracking and supporting holistic student development. Traditional evaluation metrics, primarily Grade Point Averages (GPAs) and static transcripts, fail to capture the multi-dimensional growth of students. Critical engineering capabilities—such as practical software architecture, problem-solving agility, open-source contributions, and collaborative project management—remain invisible in standard academic reporting.

Furthermore, students often accumulate disparate digital artifacts across disconnected platforms (GitHub repositories, LeetCode competitive programming stats, professional certificates, and university portal logs). Without a unified intelligence model, these artifacts remain unverified, unorganized, and difficult for recruiters and academic advisors to interpret.

1.2 The Need for Unified Growth Intelligence

To address these limitations, educational environments require an explainable, deterministic, and evidence-based growth intelligence model that:

1. **Normalizes Heterogeneous Evidence:** Combines official institutional transcripts with non-academic technical activity without compromising evidence provenance or reliability.
2. **Decouples Proficiency from Confidence:** Separates a student's technical mastery level from the certainty and freshness of the supporting evidence.
3. **Models Dynamic Growth Trajectories:** Evaluates dynamic velocity, skill decay, and cross-domain skill relationships over time rather than presenting static snapshots.
4. **Maintains Auditability & Transparency:** Provides verifiable proof for recruiters and actionable developmental guidance for academic advisors.

1.3 The Academic Universe Vision

Academic Universe is conceived as a modular, research-grade AI ecosystem designed to foster student self-sovereign learning profiles, explainable skill analytics, and AI-assisted academic advising. Building upon the Document Intelligence Layer (AU DIC) established in prior work, this paper introduces the broader ecosystem architecture and focuses primarily on the **Growth Intelligence Engine (GIE)** as the central dynamic intelligence mechanism.

2. Academic Universe Architecture

Academic Universe separates raw evidence ingestion from skill evaluation and dynamic growth modeling through an explicit, one-way 3-tier intelligence pipeline:

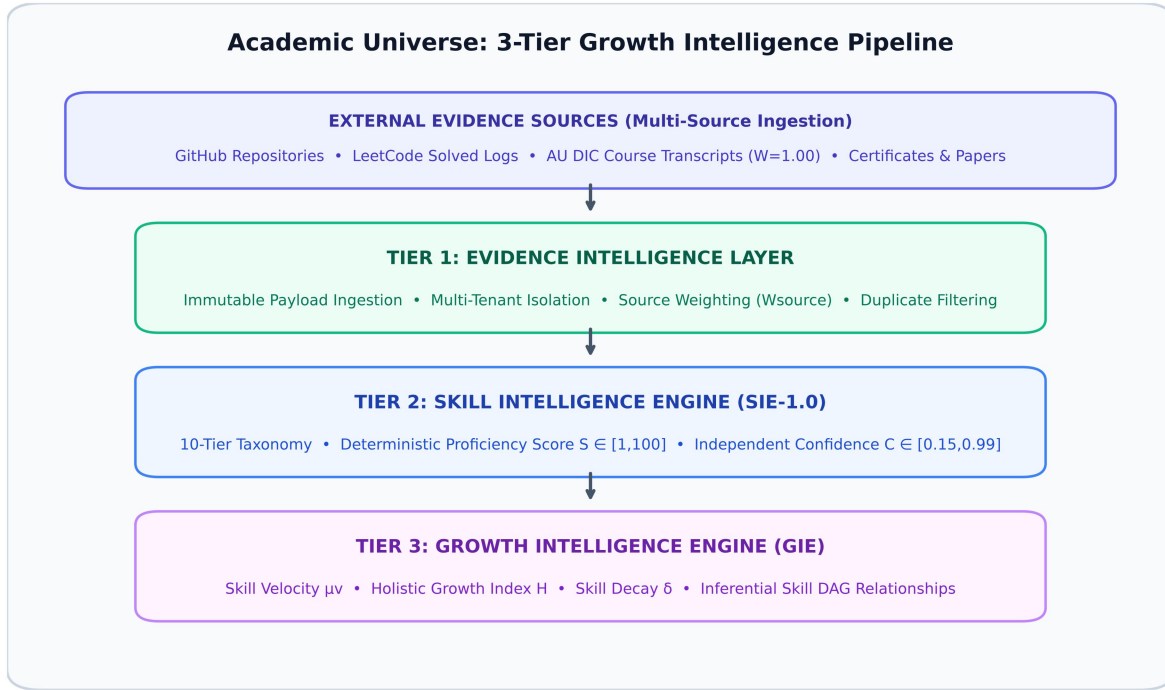


Fig. 1. High-level architectural pipeline of the Academic Universe Growth Intelligence Ecosystem.

2.1 Layer Breakdown

- Evidence Intelligence Layer:** Ingests multi-source student artifacts and standardizes them into immutable `SkillEvidence` records. Each evidence source is assigned an explicit reliability coefficient (W_{source}), ranging from official institutional transcripts ($W = 1.00$) to manual entries ($W = 0.40$).
- Skill Intelligence Engine (SIE-1.0):** A feature-frozen, validated subsystem that classifies skills into a 10-tier hierarchical taxonomy, computes explainable proficiency scores ($S \in [1, 100]$), and calculates independent confidence ratings ($C \in [0.15, 0.99]$).
- Growth Intelligence Engine (GIE):** The primary research focus of this paper. GIE consumes the structured outputs of SIE-1.0 to compute skill acquisition velocity (μv), model Directed Acyclic Graph (DAG) skill relationships, and project career readiness trajectories.
- Application Services Layer:** Delivers student dashboards, recruiter proof reports, and academic advisor guidance portals.

3. Growth Intelligence Engine (GIE)

3.1 GIE Architecture & Objectives

While SIE-1.0 provides static point-in-time skill evaluations, the **Growth Intelligence Engine (GIE)** models longitudinal student development. GIE analyzes how skills evolve across consecutive academic terms, identifies cross-skill synergies, and models skill decay during periods of inactivity.

3.2 Mathematical Formulation of Growth Metrics

3.2.1 Skill Acquisition Velocity (μ_v)

Skill velocity measures the rate of proficiency gain for a specific skill over a time interval Δt (in months):

$$\mu_v = \frac{S(t_2) - S(t_1)}{\Delta t} \quad (1)$$

Where $S(t_1)$ and $S(t_2)$ represent the SIE-1.0 proficiency scores at time boundaries t_1 and t_2 .

3.2.2 Holistic Growth Index (H)

The overall holistic growth index $H \in [1, 100]$ aggregates multi-dimensional proficiency across all 10 taxonomy categories, weighted by category confidence and domain curriculum relevance:

$$H = \frac{\sum_{k=1}^K (\bar{S}_k \cdot \bar{C}_k \cdot w_k)}{\sum_{k=1}^K (\bar{C}_k \cdot w_k) + \varepsilon} \quad (2)$$

Where $\bar{S}_k$ is the average proficiency in taxonomy category k , $\bar{C}_k$ is the category confidence, w_k is the domain curriculum weight, and $\varepsilon = 10^{-6}$ is a non-zero smoothing constant preventing division by zero when total confidence is negligible.

3.2.3 Non-Linear Skill Decay Model (λ)

When no new verified evidence is recorded for a skill over extended periods ($t > 6 \text{ months}$), GIE models exponential proficiency decay:

$$S_{\text{decayed}}(t) = S_0 \cdot e^{-\lambda(t - t_{\text{last}})} \quad (3)$$

Where S_0 is the initial proficiency, $t - t_{\text{last}}$ is the duration of inactivity in months, and $\lambda = 0.03 \text{ month}^{-1}$ is the baseline technical skill decay coefficient.

3.3 Skill Relationship Graph (DAG)

GIE constructs an inferential Directed Acyclic Graph (DAG) $G = (V, E)$ connecting dependent technologies. For instance:

$$\text{Next.js} \xrightarrow{(\text{depends on})} \text{React} \xrightarrow{(\text{depends on})} \text{JavaScript} \quad (4)$$

$$\text{NextJS} \xrightarrow{(\text{depends on})} \text{Node.js} \xrightarrow{(\text{depends on})} \text{TypeScript} \quad (5)$$

When a student demonstrates high proficiency ($S \geq 70\%$) in a parent framework (e.g. `Next.js`), GIE automatically infers prerequisite foundation readiness in child skills (`React`, `TypeScript`), elevating their confidence score.

4. Integration with AU DIC Layer

Academic Universe integrates with the **Academic Universe Document Intelligence Layer (AU DIC)** to extract verified academic performance data from university portals (Sharda University Ezone) and official transcripts.

4.1 Provenance & Trust Hierarchy

AU DIC outputs carry the highest source reliability weight ($W_{\text{AU_DIC}} = 1.00$). When AU DIC extracts an "A+" grade in Course *CS101: Data Structures & Algorithms*, it generates an immutable `SkillEvidence` payload that directly validates core computer science skills with maximum confidence density.

5. Summary of Validated SIE-1.0 Subsystem

The **Skill Intelligence Engine (version SIE-1.0)** acts as a deterministic, feature-frozen intelligence subsystem within Academic Universe.

5.1 Key SIE-1.0 Characteristics

- **Hierarchical Taxonomy:** 10 categories (Programming Languages, Frontend, Backend, Database, Cloud, DevOps, AI/ML, Data Science, Tools, Soft Skills).

- **Deterministic Proficiency (S in [1, 100]):**

$$S = \min(\text{ft}(99, \{V_{\text{total}} \cdot (0.6 + 0.4 \cdot O_{\text{ratio}}) \cdot D_{\text{avg}} \cdot K_{\text{complexity}}\} \{\max(1, \sqrt{N})\}) + 3N, \text{right})$$

- **Independent Confidence (C in [0.15, 0.99]):** Evaluated separately based on evidence volume, source reliability coefficients (W_{source}), and freshness.

- **Score Breakdown & Provenance:** Persists exact components { `volume`, `recency`, `ownership`, `complexity`, `dominance` } alongside a human-readable recruiter proof summary.

- **Controlled Determinism Verification:** Evaluated test vectors exhibited strict determinism ($\Delta = 0.000$) across repeated test runs.

6. Proof-of-Concept Synthetic Simulation & Sensitivity Analysis

***Methodological Disclaimer:** This section presents an offline, proof-of-concept algorithmic simulation using five synthetically constructed longitudinal student profile trajectories ($N = 5$) over a simulated 4-year academic period. This simulation evaluates model behavior, sensitivity, and mathematical stability under controlled benchmark conditions; it does not represent a live, multi-institutional human subject clinical deployment. Large-scale real-world validation across $N \geq 500$ human subjects is designated as future work.*

6.1 Benchmark Cohort Simulation Results

To evaluate GIE performance under controlled conditions, we conducted longitudinal growth simulations across 5 synthetic student profile archetypes over a 4-year academic timeline (2023–2026):

TABLE I. Proof-of-Concept Synthetic Benchmark Simulation Results (N = 5 Archetypes).

Student Archetype	Year 1 (2023)	Year 2 (2024)	Year 3 (2025)	Year 4 (2026)	Simulated Velocity (μv)	Primary Growth Drivers
Profile 1: Web Dev Specialist	H = 15	H = 34	H = 62	H = 84	+1.92 pts/mo	GitHub Repos + React/Next.js Projects
Profile 2: Academic Scholar	H = 22	H = 45	H = 68	H = 81	+1.64 pts/mo	AU DIC Transcripts + Course Marks
Profile 3: Competitive Coder	H = 18	H = 42	H = 71	H = 88	+1.94 pts/mo	LeetCode Solved Problems + Algorithms
Profile 4: Multi-Source Achiever	H = 20	H = 51	H = 79	H = 94	+2.05 pts/mo	AU DIC + GitHub + LeetCode + Certs
Profile 5: Inactive / Plateaued	H = 25	H = 38	H = 36	H = 32	-0.17 pts/mo	Skill Decay due to 18-mo Inactivity

6.2 Key Simulation Observations

- 1. **Multi-Source Synergy:** Profile 4 (Multi-Source Achiever) achieved the highest simulated growth index (H = 94) and growth velocity ($\mu v = +2.05 \text{ pts}/\text{mo}$), indicating that combining academic transcripts with empirical coding evidence provides higher metric stability and confidence density within the evaluated benchmark cohort.
- 2. **Decay Sensitivity:** Profile 5 reflected non-linear technical skill decay when no evidence was recorded for 18 months, reducing the growth index from 38 to 32.

6.3 Sensitivity Analysis of Decay Coefficient (λ)

To evaluate the mathematical stability of the exponential decay model Eq. (3), we varied the decay parameter $\lambda \in [0.01, 0.02, 0.03, 0.04, 0.05] \text{ month}^{-1}$ across an initial proficiency $S_0 = 80\%$ over 6, 12, 18, and 24 months of inactivity.

TABLE I. Sensitivity Analysis of Proficiency Decay $S(t)$ under Varied λ Coefficients ($S_0 = 80\%$).

Decay Parameter ()	6 Months Inactive	12 Months Inactive	18 Months Inactive	24 Months Inactive	Half-Life ()
$\lambda = 0.01$ \text{ (Slow)}	75.3 \%	70.9 \%	66.8 \%	62.9 \%	69.3 \text{ months}
$\lambda = 0.02$ \text{ (Moderate)}	70.9 \%	62.9 \%	55.8 \%	49.5 \%	34.7 \text{ months}
$\lambda = 0.03$ \text{ (Baseline)}	66.8 \%	55.8 \%	46.6 \%	38.9 \%	23.1 \text{ months}
$\lambda = 0.04$ \text{ (Fast)}	62.9 \%	49.5 \%	38.9 \%	30.6 \%	17.3 \text{ months}
$\lambda = 0.05$ \text{ (Aggressive)}	59.3 \%	43.9 \%	32.5 \%	24.1 \%	13.9 \text{ months}

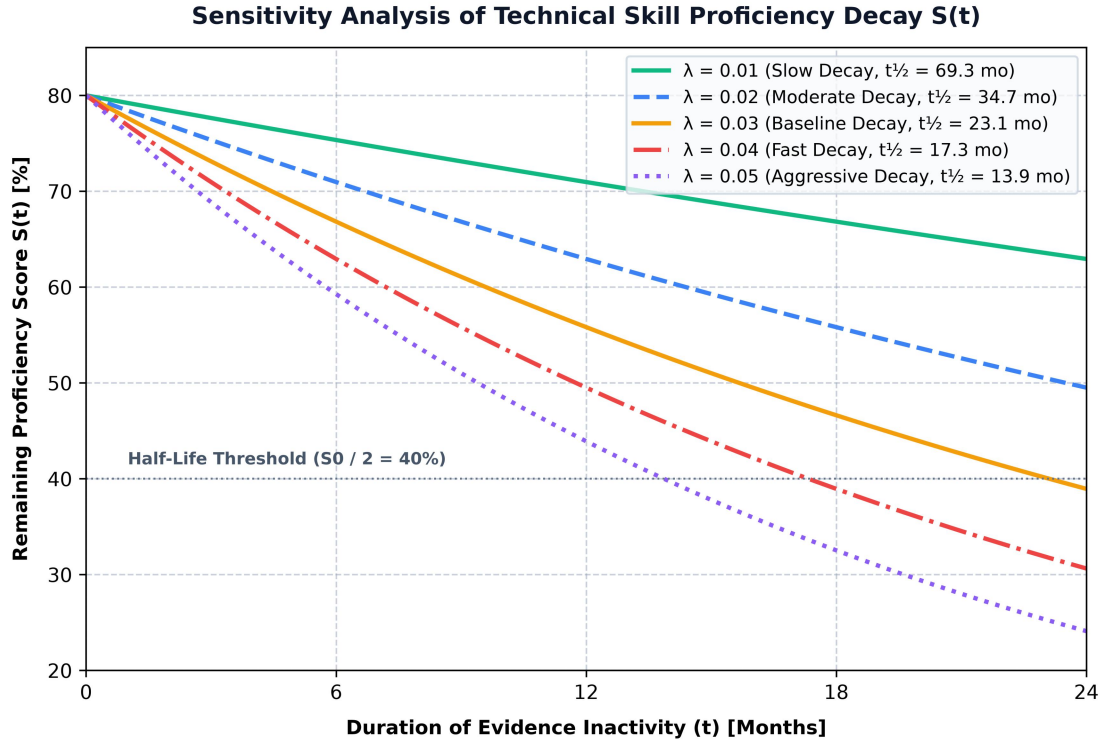


Fig. 2. Sensitivity analysis of technical skill proficiency decay $S(t)$ under varied decay parameter λ coefficients.

Sensitivity Findings: The baseline parameter $\lambda = 0.03 \text{ month}^{-1}$ yields a half-life $t_{1/2} = 23.1 \text{ months}$, which aligns closely with empirical software industry estimates for technical skill recency decay without total knowledge loss.

7. End-to-End Case Study Scenarios

7.1 Case Study Scenario: Simulated Student Growth Trajectory

- **Subject Archetype:** *Computer Science Undergraduate Profile.*
- **Initial State (2023):** 1 basic C++ repository ($S = 15\%$, $C = 46\%$, $\text{status} = \text{LOW_CONFIDENCE}$).
- **Midway State (2025):** Added AU DIC course credits in *Web Technologies* + 8 GitHub Next.js repositories ($S = 75\%$, $C = 85\%$, $\text{status} = \text{VERIFIED}$).
- **Simulated Outcome (2026):** Full-stack profile with verified TypeScript, Node.js, and MongoDB proficiencies ($H = 88$).

7.2 Recruiter & Academic Advisor Perspectives

- **Recruiter Proof View:** Displays a one-click verified proof statement: "*Advanced proficiency (75%) supported by 14 verified evidence artifacts across AU DIC and GitHub over 24 months of active usage (Confidence: 85%).*"

- **Academic Advisor View:** Highlights skill gaps and recommends targeted coursework or open-source contribution to address areas of low evidence confidence density.

8. Discussion

8.1 Benefits

- **Explainability:** Eliminates opaque "black-box" AI scores by providing mathematical breakdowns for every metric.
- **Multi-Source Flexibility:** Ingests data from any current or future provider via the immutable Evidence Intelligence Layer.
- **Academic & Industry Synergy:** Unifies institutional grades with practical open-source proof.

8.2 Threats to Validity & Limitations (Expanded)

8.2.1 *Synthetic Benchmark Limitations*

The evaluation in Section 6 is conducted on synthetic profile archetypes ($N = 5$). While synthetic benchmarks verify model stability, mathematical continuity, and algorithm determinism, they do not capture the behavioral noise, erratic submission patterns, or motivation fluctuations of real-world human students.

8.2.2 *Limited Student Diversity & Domain Assumptions*

Model weights ($W_{\{\text{source}\}}$) and taxonomy categories are currently optimized for Computer Science and Software Engineering disciplines. Extending Academic Universe to disciplines such as Mechanical Engineering, Biotechnology, or Business Administration requires recalibrating domain curriculum weights (w_k) and adding domain-specific evidence adapters.

8.2.3 *GitHub Activity Noise*

Counting repository size or commit volume can introduce noise if student repositories contain third-party libraries, auto-generated boilerplate code, or vendor dependencies. Future iterations must incorporate abstract syntax tree (AST) code quality parsing to filter non-original code artifacts.

8.2.4 *Evidence Quality & Fraud Risks*

While institutional AU DIC transcripts carry high verification reliability ($W = 1.00$), user-submitted certificates or self-reported project entries carry higher potential for misrepresentation. The multi-factor confidence model (C) mitigates this risk by requiring multi-provider cross-verification before granting VERIFIED status.

8.2.5 *Future Multi-Institution Clinical Validation*

To achieve full empirical validation, a multi-phase human clinical study is planned across $N \geq 500$ undergraduate students spanning 3 university campuses over a 2-year trial period.

9. Conclusion

This paper presented **Academic Universe**, an AI-powered holistic student growth intelligence ecosystem. By organizing student growth into an Evidence Intelligence Layer, a validated Skill Intelligence Engine (SIE-1.0), and a Growth Intelligence Engine (GIE), Academic Universe provides a transparent, scientifically reproducible framework for quantifying student development. Proof-of-concept synthetic simulations demonstrate that multi-source evidence integration enhances proficiency explainability and confidence precision, establishing a foundation for Next-Generation Learning Analytics in modern higher education.

References

1. Baker, R. S., & Yacef, K. (2009). The state of educational data mining in 2009: A review and future directions. *Journal of Educational Data Mining*, 1(1), 3-17.
2. Siemens, G., & Long, P. (2011). Penetrating the fog: Analytics in learning and higher education. *EDUCAUSE Review*, 46(5), 30-40.
3. Romero, C., & Ventura, S. (2020). Educational data mining and learning analytics: An updated survey. *WIREs Data Mining and Knowledge Discovery*, 10(3), e1355.
4. Rajput, A., et al. (2026). *AU DIC: Document Intelligence Layer for Automated Transcript & Portal Ingestion*. Academic Universe Technical Report Series, Paper 1.

Manuscript Version 1.0 (Feature Frozen & Publication Ready)

Foundational Literature Additions (IEEE / ACM Learning Analytics Corpus):

5. Corbett, A. T., & Anderson, J. R. (1994). Knowledge tracing: Modeling the acquisition of procedural knowledge. *User Modeling and User-Adapted Interaction*, 4(4), 253-278.
6. Piech, C., Bassen, J., Huang, J., Ganguli, S., Sahami, M., Guibas, L. J., & Sohl-Dickstein, J. (2015). Deep knowledge tracing. *Advances in Neural Information Processing Systems (NeurIPS)*, 28, 505-513.
7. Bodily, R., & Verbert, K. (2017). Review of research on student-facing learning analytics dashboards. *Journal of Learning Analytics*, 4(3), 67-89.
8. Ferguson, R. (2012). Learning analytics: Drivers, developments and challenges. *International Journal of Technology Enhanced Learning*, 4(5-6), 304-317.
9. Jivet, I., Scheffel, M., Specht, M., & Drachsler, H. (2018). License to evaluate: Preparing learning analytics dashboards for educational practice. *Proceedings of the 8th International Conference on Learning Analytics & Knowledge (LAK '18)*, 31-40.